



SINTEF

JutulDarcy.jl

A fully differentiable high-performance reservoir simulator based on automatic differentiation

Olav Møyner, SINTEF Digital, Applied computational sciences

ECMOR 2024, September 3rd, 2024



SINTEF

Motivation: Why a new code?

SINTEF is a contract research institute with many simultaneous projects, software is used in many projects with different needs and requires rapid turnaround

MRST is written in MATLAB/Octave and uses automatic differentiation (AD)

- Code was primarily developed for reservoir simulation, but is used for many projects
- Examples: Batteries, fuel cells, electrolysis, CO2 capture processes, ...
- Issues with MATLAB for new applications:
 - Poor performance on hard-to-vectorize and small models
 - Commercial license required
 - Difficult to deploy for clients
 - Most universities have transitioned to Python (with some Julia)

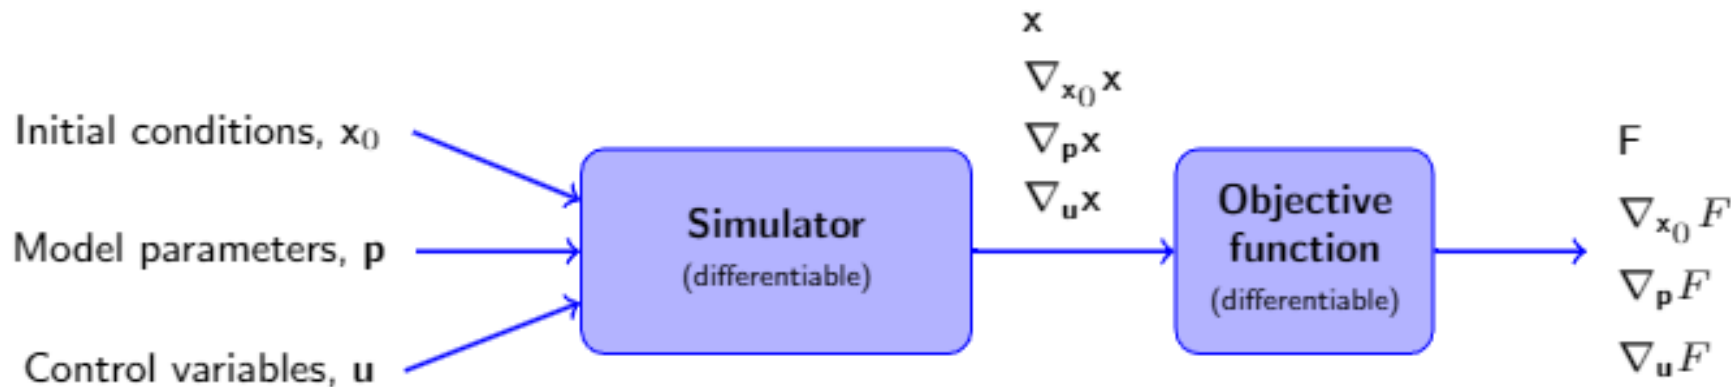
Most new models are heterogenous "multi-physics" models you do not have a single large component that can be optimized for performance (no big "reservoir" that is biggest assembly cost)



SINTEF

Motivation: Reservoir simulation

- Reservoir simulators are typically used with standard input files and standard output
- For some workflows this is a bad fit – want to "script" the simulator
- A fully differentiable simulator is not useful if it is constrained to a hard-to-build C++ code or derivatives are only accessible via IO



Taken from Odd Andersen's presentation yesterday



SINTEF

Motivation: Types of codes

Research efficiency is the time taken to arrive at a conclusion

= modelling time + simulation time + paper/report writing time

Type of code	Typical language	Simulation performance	Ease of setup	Ease of modification
HPC code	Fortran, C++	Really fast – and scales! 😊	A few days fighting CMake 😐	~O(PhD duration) 😞
Commercial code	You get a binary...	Not great, not terrible 😐	Quick and costly 😐	Impossible 💀
Flexible code	Python, MATLAB	Slow – hard to scale 😞	Quick and free 😊	Easy 😊

Motivation: Software engineering

- High level, flexible codes eventually hit a performance wall
 - You start writing extensions in C, C++ or Fortran
 - Your nice self-contained code bundle now requires a build system and careful memory management
- Low-level, high-performance codes eventually hit a flexibility wall
 - You start adding a high-level Python or MATLAB API to control your simulator
 - Difficult to translate a HPC simulator memory and execution model to a high-level API that is useful

High risk of getting the worst of both worlds!

Machine learning: Libraries have a high-level Python layer and a high-performance layer, with a **substantial engineering effort** to make deployment and install easy

Something to ponder: Would there be less papers presented on machine learning if reservoir simulators were as easy to use and as accessible as for example PyTorch?



SINTEF

Motivation: Extending models with new effects

Two physical models

Model A with linearization

$$\mathbf{R}_a(\mathbf{x}_a) = \mathbf{0}, \quad -\mathbf{J}_{aa}\Delta\mathbf{x}_a = -\mathbf{R}_a$$

Model B with linearization

$$\mathbf{R}_b(\mathbf{x}_b) = \mathbf{0}, \quad -\mathbf{J}_{bb}\Delta\mathbf{x}_b = -\mathbf{R}_b$$

Combined model

$$R(\mathbf{x}_a, \mathbf{x}_b) = \begin{bmatrix} R_a(\mathbf{x}_a, \mathbf{x}_b) \\ R_b(\mathbf{x}_b, \mathbf{x}_a) \end{bmatrix} = \mathbf{0}$$

\mathbf{J}_{aa}	\mathbf{J}_{ab}	\mathbf{x}_a	\mathbf{R}_a
\mathbf{J}_{ba}	\mathbf{J}_{bb}	\mathbf{x}_b	\mathbf{R}_b

$$=$$

Models of interest consist of *many* sub-models. Need *fast* automatic differentiation!



SINTEF

Motivation: Building blocks of the adjoint method

$$\mathbf{R}_i(\mathbf{x}_i(\mathbf{p}), \mathbf{x}_{i-1}(\mathbf{p}), \mathbf{p}, \mathbf{f}_i) = 0, \quad \forall i \in \{1, \dots, n\}$$

$$\lambda_i = - \left(\frac{\partial \mathbf{R}_i}{\partial \mathbf{x}_i} \right)^T \left(\frac{\partial O_i}{\partial \mathbf{x}_i} + \frac{\partial \mathbf{R}_{i+1}}{\partial \mathbf{x}_i} \lambda_{i+1} \right) \quad \frac{dJ_\lambda}{d\mathbf{p}} = \sum_{i=1}^n \left[\frac{\partial O_i}{\partial \mathbf{p}} + \frac{\partial \mathbf{R}_i}{\partial \mathbf{p}} \lambda_i \right]^T = \frac{dO}{d\mathbf{p}}$$

From forward simulation

$$\left(\frac{\partial \mathbf{R}_i}{\partial \mathbf{x}_i} \right)^T \quad \text{Linear solve}$$

$$\left(\frac{\partial \mathbf{R}_{i+1}}{\partial \mathbf{x}_i} \right) \quad \text{Matrix-vector product}$$

New shape, but uses internals only

$$\left(\frac{\partial \mathbf{R}_i}{\partial \mathbf{p}} \right)^T \quad \text{Matrix-vector product}$$

User provided code!

$$\left(\frac{\partial O_i}{\partial \mathbf{x}} \right)^T \quad \text{Element-wise addition}$$

$$\left(\frac{\partial O_i}{\partial \mathbf{p}} \right)^T \quad \text{Element-wise addition}$$

Fully differentiable simulators

Goal: A fully differentiable, scriptable, high-performance porous media simulator


1. Does not need to be *the* fastest, but must be *competitive*
2. Does not need to run *all* models, but must *run relevant models*
3. Try to build on experiences from AD in AD-GPRS, MRST and OPM Flow

Many alternatives:

- Python with accelerators: Numba/JAX and in the future Mojo?
- Bolt something onto machine learning libraries (PyTorch/TensorFlow)?
- Modern C++/Fortran?
- Adapt existing codes?
- New languages: Chapel, Rust, Carbon, Julia?

Main considerations:

- Native support for automatic differentiation for coupled models
- Interactive development without big compilation toolchain
- Maturity (per 2020) and available building blocks:





Jutul

- Julia package for implicit solves of coupled models
- Automatic differentiation of discrete equations
- Design inspired by use of AD in MRST and OPM Flow
- Robust Newton solvers (chopping, relaxation, variable scaling, absolute and relative change limits)
- Interactive visualization

www.github.com/sintefmath/Jutul.jl

Reservoir simulation



[JutulDarcy.jl](#) is a high performance Darcy flow simulator and the main demonstrator application for Jutul. See also [JutulDarcyRules.jl](#) for use in differentiable workflows involving CO2 storage.

Battery simulation



[BattMo.jl](#) is a battery simulator that implements a subset of the MATLAB-based [BattMo](#) toolbox in Julia for improved performance.

Carbon capture

Jutul.jl powers a simulator that implements vacuum swing adsorption and direct air capture processes for the capture of CO2. This application is currently not public.



JutulDarcy demonstrator

Main demonstrator is a porous media simulator:

- Immiscible, thermal, black-oil and equation-of-state or K-value compositional flow
- Industry standard input – or write scripts where you define grid, properties and schedule yourself
- *Fully differentiable*: Gradients with respect to any declared parameter using adjoint method
- MPI parallel with BoomerAMG and threads
- MIT licensed, open source at www.github.com/sintefmath/JutulDarcy.jl
- Installation on any OS (Julia 1.8+):
`using Pkg; Pkg.add("JutulDarcy")`



Compact implementation

JutulDarcy: 9000 lines of code (loc)

- **Immiscible flow base 1000 loc**
- **Black-oil specialization 1100 loc**
- **Multisegment wells with advanced control logic 1430 loc**
- **Compositional 705 loc**



SINTEF

Correspondence between equations and code

- All equation terms are written as functions without type annotations or templates
- Dispatch is used to distinguish different types of physics and discretizations
- Equations are easily recognizable in code

$$\frac{\partial}{\partial t} [\phi \rho_o^s (b_o S_o + R_s b_g S_g)] + \rho_o^s \nabla \cdot (b_o \vec{v}_o + R_s b_o \vec{v}_o) - q_o = 0,$$
$$\frac{\partial}{\partial t} [\phi \rho_g^s (b_g S_g + R_v b_o S_o)] + \rho_g^s \nabla \cdot (b_g \vec{v}_g + R_v b_g \vec{v}_o) - q_g = 0$$

```
1 function update_blackoil_mass!(M, pv, b, S, Rs, Rv, rhoS, sys)
2     Φ = pv[cell]                                # pore-volume
3     l, v = phase_indices(sys)                    # 1, 2 or 2, 1
4     bO, bG = b[l, cell], b[v, cell]              # get b-factors
5     sO, sG = S[l, cell], S[v, cell]              # get saturations
6     M[l] = Φ*rhoS[l]*(bO*sO + bG*sG*Rv[cell])    # oil component mass
7     M[v] = Φ*rhoS[v]*(bG*sG + bO*sO*Rs[cell])    # gas component mass
8     return M
9 end
```




SINTEF

Example: Black-oil flux implementation

```
# Oil
f_b1 = cell -> b_mob(b, kr, μ, l, cell)
λb_l = upwind(upw, f_b1, ψ_l)
q_l = rhoS[l]*λb_l*ψ_l
# Gas mobility
f_bv = cell -> b_mob(b, kr, μ, v, cell)
λb_v = upwind(upw, f_bv, ψ_v)
# Rs (solute gas) upwinded by liquid potential
f_rs = cell -> @inbounds Rs[cell]
rs = upwind(upw, f_rs, ψ_l)
# Final flux = gas phase flux + gas-in-oil flux
q_v = (λb_v*ψ_v + rs*λb_l*ψ_l)*rhoS[v]
```


```
function b_mob(b, kr, μ, ph, c)
    λ = kr[ph, c]/μ[ph, c]
    b_f = b[ph, c]
    return λ*b_f
end
```

```
function upwind(upw::SPU, F, q)
    flag = q >= 0
    if flag
        up = upw.right
    else
        up = upw.left
    end
    return F(up)
end
```



SINTEF

Fast automatic differentiation



Jutul.jl uses a tailored approach automatic differentiation that:

- Supports complex sparsity patterns (unstructured grids and complex wells)
- Allows for equations and variables on different entities (cells, faces, wells, ...)
- Key ideas:
 - Detect sparsity pattern of equations at start of simulation
 - Allocate dense memory for storage of dual numbers, in the order of evaluation
 - Manage properties in a precomputed dependency graph
 - Avoid (significant) heap allocation during simulation and make GC happy

**General AD for any discrete equation
(Wells, facilities, coupling terms)**



**General finite-volume scheme
(AvgMPFA, NTPFA...)**



Two-point flux scheme

Specialization and performance






Case	Model	No. Cells	No. Wells	No. Assembly	Runtime	AD + props + assembly
SPE1	Black oil, 3-phase	300	2	521 (0.26 ms)	0.3 s	24 %
SPE9	Black-oil, 3-phase	9 000	26	403 (6.5 ms)	6.1 s	26 %
Egg	Water-oil, 2-phase	18 553	12	706 (4.9 ms)	15.8 s	11 %
Norne	Black-oil, 3-phase	44 431	36	2597 (41.5 ms)	260 s	18 %
Olympus	Water-oil, 2-phase	192 749	18	807 (62.3 ms)	162 s	18 %
Sleipner	Water-gas, 2-phase (CO2)	1 986 176	1	1965 (704 ms)	4 519 s	18 %
A	7 components, 3-phase	60 000	2	2097 (397 ms)	1071 s	43 %
B	Water-oil, 2-phase	200 000	100+	2299 (65 ms)	572 s	14 %
C	Black-oil, 3-phase	150 000	250+	18802 (132 ms)	10 205 s	16 %
D	Black-oil, 3-phase	1 250 000	100+	7259 (837 ms)	41 129 s	9.4 %

All cases run with –O3 single thread and default options on Ryzen 9 16 core CPU

Open models: Check against your favorite simulator!

References: 10.2118/9723-PA, 10.3997/2214-4609.201802246, 10.2118/182679-PA 10.1002/gdj3.21, 10.2118/127538-MS, 10.2118/29110-MS, 10.2118/72469-PA



Validation: OLYMPUS 1





SINTEF

Verification: SPE9


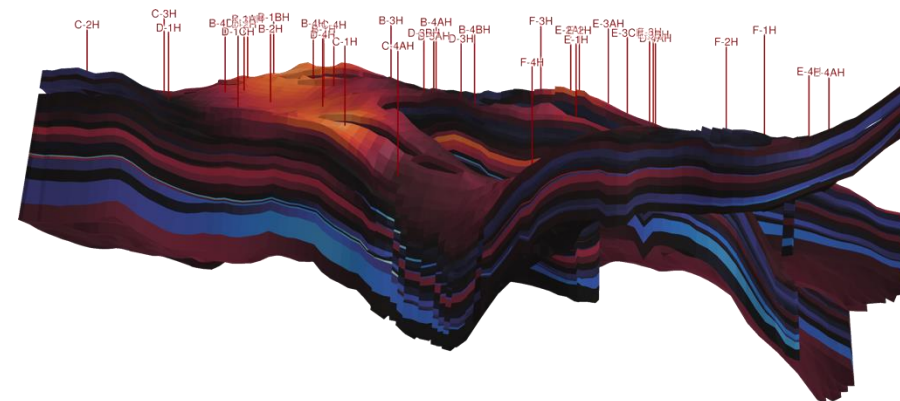


— PROD25	— PROD16	— PROD20
— PROD10	— PROD12	— PROD22
— PROD14	— PROD6	— PROD7
— PROD9	— PROD24	— PROD3
— PROD23	— PROD5	— PROD11
— PROD17	— PROD4	— PROD15
— PROD2	— PROD19	— PROD21
— PROD13	— PROD8	— PROD26
— PROD18		



SINTEF

Validation: Norne




D-2H	D-4H	B-4H
B-4AH	B-2H	E-1H
B-4BH	D-3AH	D-3BH
E-3H	K-3H	E-4AH
D-1H	B-1AH	E-3CH
E-4H	E-2H	E-2AH
B-1BH	B-4DH	D-3H
E-3AH	D-4AH	B-3H
E-3BH	D-1CH	B-1H



SINTEF

A toy example: Quarter-five with barriers

- Standard conceptual test case
- Gas is injected in bottom left and liquid is produced in upper right
- Layers with different rock types impede flow
- Viscous driven, no gravity
- How easy is it to modify a pre-existing model?



(a) Permeability (darcy)



SINTEF

Adding properties and parameters

Define new relative permeability definition and evaluator function:

```
1 import JutulDarcy: AbstractRelativePermeabilities
2 struct MyKr <: AbstractRelativePermeabilities end
3 @jutul_secondary function update_my_kr!(vals, def::MyKr, model, Saturations,
4     KrExponents, cells_to_update)
5     for c in cells_to_update
6         for ph in axes(vals, 1)
7             S_α = max(Saturations[ph, c], 0.0)
8             n_α = KrExponents[ph, c]
9             vals[ph, c] = S_α^n_α
10         end
11     end
12 end
```

$$k_{rw} = S_{\alpha}^n, k_{rg} = S_{\alpha}^m$$

Exponents n, m vary spatially

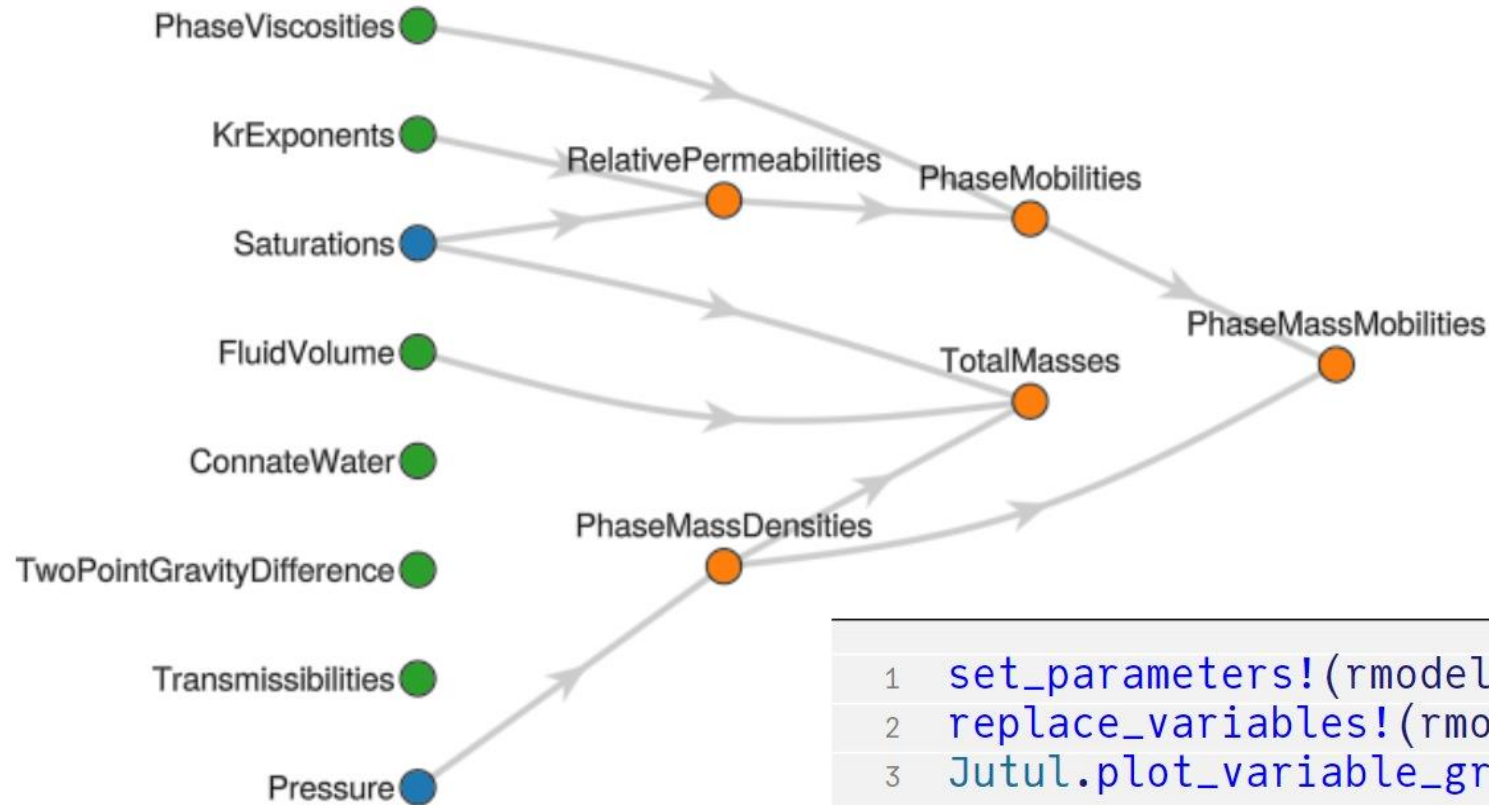
Add a new parameter for Corey exponent:

```
1 import JutulDarcy: PhaseVariables
2 struct MyKrExp <: PhaseVariables end
3 Jutul.default_value(model, ::MyKrExp) = 2.0
```



SINTEF

Dependency graph autogenerated!



```
1 set_parameters!(rmodel, KrExponents = MyKrExp())
2 replace_variables!(rmodel, RelativePermeabilities = MyKr())
3 Jutul.plot_variable_graph(rmodel)
```


● Primary variable ● Secondary variable ● Parameter



SINTEF

Running the "new" simulator

- We set the values of the new parameter to vary with rock types
- Running a simulation immediately recompiles only small parts of the code
- Performance of new function is excellent – no difference in runtime after adding
- The inserted function could also be a neural network or a call to an external package



(b) Progression of gas saturation



SINTEF

Let us calculate some gradients!

- Gradient calculations: Let us define the producer gas rate as objective

```
1 pv = pore_volume(model, prm)
2 total_time = sum(dt)
3 inj_rate = sum(pv)/total_time
4 import JutulDarcy: compute_well_qoi
5 function objective_function(model, state, Δt, step_i, forces)
6     T = SurfaceGasRateTarget
7     grat = compute_well_qoi(model, state, forces, :Producer, T)
8     return Δt*grat/(inj_rate*total_time)
9 end
```


- Call high level interface to get gradients with respect to *input* parameters (perm, poro, geometry) rather than *numerical* parameters (transmissibilities and pore-volumes)

```
1 import JutulDarcy: reservoir_sensitivities
2 grad = reservoir_sensitivities(case, result, objective_function)
```



SINTEF

Examples of sensitivities from one adjoint solve




Cell depth




Permeability




Porosity



Liquid Corey exponent



Gas Corey exponent




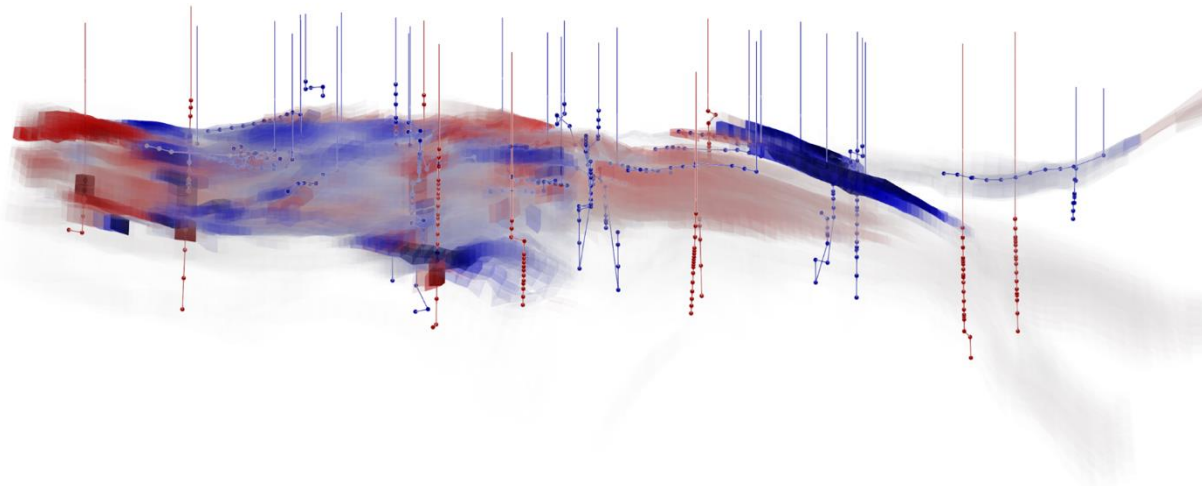
Liquid viscosity



SINTEF

Gradients: Norne gas production

```
total_time = sum(case.dt) # Normalize objective by total time
import JutulDarcy: compute_well_qoi, reservoir_sensitivities
function objective(model, state, Δt, step_i, forces)
    grat = 0.0
    T = SurfaceGasRateTarget # Get out gas rate
    for w in prod
        grat += compute_well_qoi(model, state, forces, w, T)
    end
    return Δt*grat/total_time # Total gas production, normalized by time
end
```





SINTEF

134k variables: 1 million numerical parameters

Parameter	Count	Type	Note
Transmissibilities	132693×1	Faces	Discretizes potential difference
Gravity difference $g\Delta z$	132693×1	Faces	Discretizes bouyancy
Connate water	44417×1	Cells	Used in evaluation of k_r
Water k_r scalars	44417×4	Cells	End-points and maximum k_r
Oil-water k_r scalars	44417×4	Cells	–
Oil-gas k_r scalars	44417×4	Cells	–
Gas k_r scalars	44417×4	Cells	–
Pore-volume	44417×1	Cells	Pore space available to flow
Well indices	503×1	Perforations	Connection well and reservoir
Perforation Δz	503×1	Perforations	Depth from well to perforation
Top node volume	36×1	Well	Volume for standard well



SINTEF

Input variables for Norne

Parameter	Count	Type	Impact
Areas	132693×1	Faces	Transmissibilities
Face centroids	132693×3	Faces	Transmissibilities
Normals	132693×3	Faces	Transmissibilities
Transmissibility multiplier	132693×1	Faces	Transmissibilities
Permeability	44417×3	Cells	Transmissibilities and well indices
Porosity	44417×1	Cells	Pore-volume
Net-to-gross	44417×1	Cells	Pore-volume, transmissibilities, well indices
Water k_r scalars	44417×4	Cells	Direct copy
Oil-water k_r scalars	44417×4	Cells	Direct copy
Oil-gas k_r scalars	44417×4	Cells	Direct copy
Gas k_r scalars	44417×4	Cells	Direct copy
Volumes	44417×1	Cells	Pore-volumes
Cell centroids	44417×3	Cells	$g\nabla z$, well perforation depth difference and transmissibilities



SINTEF

Concluding remarks

- Automatic differentiation (AD) has a reputation as slow
 - Careful application can give excellent performance and unparalleled flexibility
- Adjoint, fast AD and a systematic approach to parameters allows immediate understanding of what parameters impact a given objective function
- Ideas are not Julia specific - can easily be adapted to other codes
- Julia largely delivers on the promise of both scripting flexibility and performance
 - Some issues remain, like poor debugging support and latency when recompilation is required
 - Performance is excellent and multiple dispatch is very powerful
 - Package manager is great – and a good answer to build system despair and dependency hell

```
using Pkg; Pkg.add(["JutulDarcy", "MPI", "HYPRE", "GLMakie"])
```

Thank you for your attention!

<https://github.com/sintefmath/JutulDarcy.jl>